

# BDAT 1008 - 04 Data Collection and Curation Assignment 2

Shubham Chawla 200493036

Question 1 A sample schema for orders\_data is given below. Import the data. How many unique cities are there?

In [73]:

```
# Importing Necessary Libraries
from pyspark.sql.types import *
from pyspark.sql import functions as F
from pyspark.sql import DataFrameWriter as W
from math import radians, cos, sin, asin, sqrt

# Spark Session
from pyspark.sql import SparkSession
spark = SparkSession.builder.getOrCreate()
spark.conf.set('spark.sql.shuffle.partitions', 6)
spark.conf.set('num-executors', 16)
spark
```

Out[73]: **SparkSession - in-memory**

**SparkContext**

[Spark UI](#)

<b>Version</b>	v3.2.1
<b>Master</b>	local[*]
<b>AppName</b>	pyspark-shell

In [74]:

```
%ls orders_data
```

```
Volume in drive C is OS
Volume Serial Number is 3C52-E078

Directory of C:\Users\shubh\Desktop\1008 Data Collection & Curation\A2\orders_data

2022-07-10  05:59 PM    <DIR>          .
2022-07-17  06:58 PM    <DIR>          ..
2022-07-10  05:59 PM                14,173 orders_1.csv
2022-07-10  05:59 PM                14,413 orders_10.csv
2022-07-10  05:59 PM                14,392 orders_100.csv
2022-07-10  05:59 PM                14,390 orders_11.csv
2022-07-10  05:59 PM                14,432 orders_12.csv
2022-07-10  05:59 PM                14,287 orders_13.csv
2022-07-10  05:59 PM                14,354 orders_14.csv
2022-07-10  05:59 PM                14,465 orders_15.csv
2022-07-10  05:59 PM                14,442 orders_16.csv
2022-07-10  05:59 PM                14,592 orders_17.csv
2022-07-10  05:59 PM                14,439 orders_18.csv
2022-07-10  05:59 PM                14,455 orders_19.csv
2022-07-10  05:59 PM                14,485 orders_2.csv
2022-07-10  05:59 PM                14,493 orders_20.csv
2022-07-10  05:59 PM                14,461 orders_21.csv
2022-07-10  05:59 PM                14,449 orders_22.csv
2022-07-10  05:59 PM                14,518 orders_23.csv
2022-07-10  05:59 PM                14,416 orders_24.csv
2022-07-10  05:59 PM                14,464 orders_25.csv
2022-07-10  05:59 PM                14,492 orders_26.csv
2022-07-10  05:59 PM                14,432 orders_27.csv
2022-07-10  05:59 PM                14,391 orders_28.csv
2022-07-10  05:59 PM                14,414 orders_29.csv
2022-07-10  05:59 PM                14,460 orders_3.csv
2022-07-10  05:59 PM                14,489 orders_30.csv
2022-07-10  05:59 PM                14,488 orders_31.csv
2022-07-10  05:59 PM                14,419 orders_32.csv
2022-07-10  05:59 PM                14,451 orders_33.csv
2022-07-10  05:59 PM                14,406 orders_34.csv
2022-07-10  05:59 PM                14,378 orders_35.csv
2022-07-10  05:59 PM                14,449 orders_36.csv
2022-07-10  05:59 PM                14,406 orders_37.csv
2022-07-10  05:59 PM                14,470 orders_38.csv
2022-07-10  05:59 PM                14,460 orders_39.csv
2022-07-10  05:59 PM                14,156 orders_4.csv
2022-07-10  05:59 PM                14,565 orders_40.csv
2022-07-10  05:59 PM                14,655 orders_41.csv
2022-07-10  05:59 PM                14,438 orders_42.csv
2022-07-10  05:59 PM                14,350 orders_43.csv
2022-07-10  05:59 PM                14,477 orders_44.csv
2022-07-10  05:59 PM                14,445 orders_45.csv
2022-07-10  05:59 PM                14,414 orders_46.csv
2022-07-10  05:59 PM                14,506 orders_47.csv
2022-07-10  05:59 PM                14,524 orders_48.csv
2022-07-10  05:59 PM                14,486 orders_49.csv
```

2022-07-10	05:59 PM	14,313	orders_5.csv
2022-07-10	05:59 PM	14,427	orders_50.csv
2022-07-10	05:59 PM	14,502	orders_51.csv
2022-07-10	05:59 PM	14,419	orders_52.csv
2022-07-10	05:59 PM	14,477	orders_53.csv
2022-07-10	05:59 PM	14,427	orders_54.csv
2022-07-10	05:59 PM	14,470	orders_55.csv
2022-07-10	05:59 PM	14,420	orders_56.csv
2022-07-10	05:59 PM	14,534	orders_57.csv
2022-07-10	05:59 PM	14,442	orders_58.csv
2022-07-10	05:59 PM	14,482	orders_59.csv
2022-07-10	05:59 PM	14,308	orders_6.csv
2022-07-10	05:59 PM	14,482	orders_60.csv
2022-07-10	05:59 PM	14,495	orders_61.csv
2022-07-10	05:59 PM	14,476	orders_62.csv
2022-07-10	05:59 PM	14,441	orders_63.csv
2022-07-10	05:59 PM	14,471	orders_64.csv
2022-07-10	05:59 PM	14,450	orders_65.csv
2022-07-10	05:59 PM	14,416	orders_66.csv
2022-07-10	05:59 PM	14,384	orders_67.csv
2022-07-10	05:59 PM	14,545	orders_68.csv
2022-07-10	05:59 PM	14,407	orders_69.csv
2022-07-10	05:59 PM	14,473	orders_7.csv
2022-07-10	05:59 PM	14,423	orders_70.csv
2022-07-10	05:59 PM	14,571	orders_71.csv
2022-07-10	05:59 PM	14,433	orders_72.csv
2022-07-10	05:59 PM	14,488	orders_73.csv
2022-07-10	05:59 PM	14,480	orders_74.csv
2022-07-10	05:59 PM	14,537	orders_75.csv
2022-07-10	05:59 PM	14,465	orders_76.csv
2022-07-10	05:59 PM	14,421	orders_77.csv
2022-07-10	05:59 PM	14,332	orders_78.csv
2022-07-10	05:59 PM	14,402	orders_79.csv
2022-07-10	05:59 PM	14,465	orders_8.csv
2022-07-10	05:59 PM	14,460	orders_80.csv
2022-07-10	05:59 PM	14,457	orders_81.csv
2022-07-10	05:59 PM	14,478	orders_82.csv
2022-07-10	05:59 PM	14,419	orders_83.csv
2022-07-10	05:59 PM	14,398	orders_84.csv
2022-07-10	05:59 PM	14,429	orders_85.csv
2022-07-10	05:59 PM	14,331	orders_86.csv
2022-07-10	05:59 PM	14,510	orders_87.csv
2022-07-10	05:59 PM	14,345	orders_88.csv
2022-07-10	05:59 PM	14,364	orders_89.csv
2022-07-10	05:59 PM	14,434	orders_9.csv
2022-07-10	05:59 PM	14,393	orders_90.csv
2022-07-10	05:59 PM	14,449	orders_91.csv
2022-07-10	05:59 PM	14,496	orders_92.csv
2022-07-10	05:59 PM	14,474	orders_93.csv
2022-07-10	05:59 PM	14,465	orders_94.csv
2022-07-10	05:59 PM	14,388	orders_95.csv
2022-07-10	05:59 PM	14,369	orders_96.csv

```
2022-07-10  05:59 PM          14,482 orders_97.csv
2022-07-10  05:59 PM          14,480 orders_98.csv
2022-07-10  05:59 PM          14,424 orders_99.csv
          100 File(s)      1,444,064 bytes
          2 Dir(s)  78,929,788,928 bytes free
```

```
In [75]: # Importing the data and applying sample schema
from pyspark.sql.types import *

file_location = "orders_data\orders_*.csv"

ordersSchema = StructType([
    StructField("Order_ID", DoubleType(), True),
    StructField("Country", StringType(), True),
    StructField("Province", StringType(), True),
    StructField("City", StringType(), True),
    StructField("Latitude", DoubleType(), True),
    StructField("Longitude", DoubleType(), True),
    StructField("TimeStamp", StringType(), True),
    StructField("Sales_Volume", DoubleType(), True)])

print(type(ordersSchema))
```

<class 'pyspark.sql.types.StructType'>

```
In [76]: # creating DataFrame
data = (
    spark
        .read
        .schema(ordersSchema)
        .csv(file_location)
)

# create table for SQL analytics
data.createOrReplaceTempView("orders")

# Checking Size of pyspark dataframe
print((data.count(), len(data.columns)))

data.show(5)
```

```
(20000, 8)
+-----+-----+-----+-----+-----+-----+-----+
|Order_ID|Country|Province|      City| Latitude|Longitude|      TimeStamp|Sales_Volume|
+-----+-----+-----+-----+-----+-----+-----+
|231542.0| Canada|    AB|    Calgary|-113.9835|-113.9389|2022/04/22 08:28:49|      41.49|
|473450.0| Canada|    AB|   Edmonton|-113.4467|-113.3654|2022/04/22 05:04:24|      48.52|
|376604.0| Canada|    AB| Medicine Hat|-110.5798|-110.4884|2022/04/22 18:14:14|      60.79|
|440105.0| Canada|    AB|Sherwood Park|-113.2427| -113.242|2022/04/22 11:27:20|      77.81|
|483058.0| Canada|    AB|    Beaumont|-113.3783|-113.2894|2022/04/22 21:04:24|      12.06|
+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

```
In [77]: # Checking for duplicate order ids
if data.count() > data.dropDuplicates(['Order_ID']).count():
    print('Data has {} duplicate OrderIDs'.format(data.count() - data.dropDuplicates(['Order_ID']).count()))
```

Data has 648 duplicate OrderIDs

```
In [78]: # Removing Duplicate OrderIDs
duplicates_removed_df = data.dropDuplicates(['Order_ID'])
# create table for SQL analytics - duplicates removed
duplicates_removed_df.createOrReplaceTempView("orders_duplicates_removed")
# Checking Size of dataframe after removing duplicates
print((duplicates_removed_df.count(), len(duplicates_removed_df.columns)))
```

(19352, 8)

```
In [79]: # Unique Cities Count without removing duplicates
df_sql = spark.sql("SELECT count(distinct(City)) AS `Unique Cities Count` FROM orders")
df_sql.show()
```

+-----+	
Unique Cities Count	
+-----+	
	619
+-----+	

```
In [80]: # Unique Cities Count after removing duplicate Order_ID
df_sql = spark.sql("SELECT count(distinct(City)) AS `Unique Cities Count` FROM orders_duplicates_removed")
df_sql.show()
```

+-----+	
Unique Cities Count	
+-----+	
	612
+-----+	

Question 2 For each province, show the latest time of the order that was made

```
In [83]: df_sql = spark.sql("""
select Province, max(TimeStamp) AS `Latest Time of Order`
from orders_duplicates_removed
group by Province
order by Province
""")
df_sql.show()
```

+-----+	
Province Latest Time of Order	
+-----+	
	AB  2022/04/23 01:13:04
	BC  2022/04/23 01:05:14
	MB  2022/04/23 00:25:42
	NB  2022/04/23 00:47:52
	NL  2022/04/23 00:36:20
	NS  2022/04/23 00:19:17
	NT  2022/04/22 17:33:18
	ON  2022/04/23 01:18:38
	PE  2022/04/22 23:23:21
	QC  2022/04/23 01:17:12
	SK  2022/04/23 01:19:30
	YT  2022/04/22 20:48:47
+-----+	

Question 3 A point of interest (POI) is a specific point location that people take interest. Import the POI data set (POI.csv).

```
In [84]: # Importing data using inferred schema
# Location and type of the file being imported
path = 'POI.csv'
file_type = "csv"

# CSV OPTIONS
infer_schema = "true" #By using this option spark will automatically go through the csv file and infer the schema of each column.
first_row_is_header = "true"
```

```
delimiter = ","

# Importing POI.csv file
df_POI = spark.read.format(file_type) \
    .option("inferSchema",infer_schema) \
    .option("header", first_row_is_header) \
    .option("sep", delimiter) \
    .load(path)
```

In [85]: df\_POI.columns

Out[85]: ['POI\_ID', ' Latitude', 'Longitude']

In [86]: *#renaming column Latitude to remove an extra space in the beginning*  
df\_POI = df\_POI.withColumnRenamed(" Latitude","Latitude")  
df\_POI.columns

Out[86]: ['POI\_ID', 'Latitude', 'Longitude']

In [87]: *# create view for SQL analytics*  
df\_POI.createOrReplaceTempView("POI")  
df\_POI.show(5)

POI_ID	Latitude	Longitude
POI-1	45.521629	-74.085634
POI-2	53.596345	-114.465122
POI-3	44.897823	-62.987528

**Question 4** For each Order\_ID and POI pair, get the distance between the Order\_ID and the POI based on the geographic Latitude and Longitude.

In [88]: *# Joining views orders and POI (A = orders, B = POI)*  
combined\_df = spark.sql(  
 """select int(v1.Order\_ID), v2.POI\_ID, v1.Latitude AS `Lat\_A`, v1.Longitude AS `Long\_A`, v2.Latitude AS `Lat\_B`, v2.Longitude AS `Long\_B`  
 from orders\_duplicates\_removed v1, POI v2  
 """  
)  
combined\_df.show(10)

Order_ID	POI_ID	Lat_A	Long_A	Lat_B	Long_B
250410	POI-1	-113.2142	-113.1836	45.521629	-74.085634
250410	POI-2	-113.2142	-113.1836	53.596345	-114.465122
250410	POI-3	-113.2142	-113.1836	44.897823	-62.987528
286817	POI-1	-111.4565	-111.4168	45.521629	-74.085634
286817	POI-2	-111.4565	-111.4168	53.596345	-114.465122
286817	POI-3	-111.4565	-111.4168	44.897823	-62.987528
298450	POI-1	-113.3683	-113.3579	45.521629	-74.085634
298450	POI-2	-113.3683	-113.3579	53.596345	-114.465122
298450	POI-3	-113.3683	-113.3579	44.897823	-62.987528
480255	POI-1	-79.3772	-79.3679	45.521629	-74.085634

only showing top 10 rows

In [90]:

```
from IPython.display import Image
Image(url= "formula.png", width=600, height=400)
```



In [91]:

```
# Calculating distance between two points using Haversine Distance formula

import pyspark.sql.functions as F
combined_df = combined_df.withColumn("a", (
    F.pow(F.sin(F.radians(F.col("Lat_B")) - F.col("Lat_A"))) / 2, 2) +
    F.cos(F.radians(F.col("Lat_A"))) * F.cos(F.radians(F.col("Lat_B"))) *
    F.pow(F.sin(F.radians(F.col("Long_B")) - F.col("Long_A"))) / 2, 2)
```



```
)).withColumn("Distance", F.atan2(F.sqrt(F.col("a")), F.sqrt(-F.col("a") + 1)) * 12742.018) #2*uE = 2 * 6371.009

combined_df = combined_df.drop('a')
combined_df = combined_df.withColumn("Distance", F.round(combined_df["Distance"], 3))
combined_df.show()
```

Order_ID	POI_ID	Lat_A	Long_A	Lat_B	Long_B	Distance
250410	POI-1	-113.2142	-113.1836	45.521629	-74.085634	16731.229
250410	POI-2	-113.2142	-113.1836	53.596345	-114.465122	18546.88
250410	POI-3	-113.2142	-113.1836	44.897823	-62.987528	16216.34
286817	POI-1	-111.4565	-111.4168	45.521629	-74.085634	16702.623
286817	POI-2	-111.4565	-111.4168	53.596345	-114.465122	18345.494
286817	POI-3	-111.4565	-111.4168	44.897823	-62.987528	16232.481
298450	POI-1	-113.3683	-113.3579	45.521629	-74.085634	16731.879
298450	POI-2	-113.3683	-113.3579	53.596345	-114.465122	18564.407
298450	POI-3	-113.3683	-113.3579	44.897823	-62.987528	16213.075
480255	POI-1	-79.3772	-79.3679	45.521629	-74.085634	13892.397
480255	POI-2	-79.3772	-79.3679	53.596345	-114.465122	14961.485
480255	POI-3	-79.3772	-79.3679	44.897823	-62.987528	13859.726
433267	POI-1	-82.978	-82.895	45.521629	-74.085634	14296.757
433267	POI-2	-82.978	-82.895	53.596345	-114.465122	15286.753
433267	POI-3	-82.978	-82.895	44.897823	-62.987528	14261.037
200925	POI-1	-79.3595	-79.3471	45.521629	-74.085634	13890.402
200925	POI-2	-79.3595	-79.3471	53.596345	-114.465122	14959.962
200925	POI-3	-79.3595	-79.3471	44.897823	-62.987528	13857.713
346582	POI-1	-114.3696	-114.3627	45.521629	-74.085634	16736.781
346582	POI-2	-114.3696	-114.3627	53.596345	-114.465122	18676.975

only showing top 20 rows

Question 5 For each Order\_ID, identify the POI with the shortest distance. Retain only 1 record for each Order\_ID. (Check: Your end result should have the same record count as your orders\_data dataset.)

In [101...

```
# create view for SQL analytics
combined_df.createOrReplaceTempView("orders_POI")

# Finding POI with shortest distance for each Order_ID by using window function
shortest_df = spark.sql(
    """WITH cte AS (
        SELECT Order_ID, POI_ID, Distance,
               RANK() OVER ( PARTITION BY Order_ID
                             ORDER BY Distance ASC
                           ) AS r
    ) AS r
```

```
        FROM orders_POI
    )
    SELECT Order_ID, POI_ID, Distance AS `Shortest Distance`
    FROM cte
    WHERE r = 1
    ORDER BY Order_ID;
""
)
shortest_df.show()
```

Order_ID	POI_ID	Shortest Distance
200021	POI-3	16204.454
200026	POI-3	13901.091
200041	POI-3	16209.39
200099	POI-3	13889.543
200137	POI-3	13954.091
200154	POI-3	13873.513
200171	POI-3	16167.789
200183	POI-3	13188.207
200220	POI-3	13856.457
200221	POI-3	13880.355
200231	POI-3	13857.506
200242	POI-3	13877.028
200308	POI-3	16213.982
200311	POI-3	13840.551
200313	POI-3	13857.452
200335	POI-3	13621.158
200343	POI-3	13883.572
200347	POI-3	13194.087
200374	POI-3	12933.142
200398	POI-3	14923.097

only showing top 20 rows

```
In [102... # Checking size of dataframe after retaining only single record for each unique Order_ID
print((shortest_df.count(), len(shortest_df.columns)))

(19352, 3)
```

```
In [103... # Checking order_data dataset size with duplicates removed
print((duplicates_removed_df.count(), len(duplicates_removed_df.columns)))

(19352, 8)
```

Question 6 Based on #5, for each POI, get the average, standard deviation and max of the (shortest) distances, as well as the

count of the orders.

In [107...

```
# create view for SQL analytics
shortest_df.createOrReplaceTempView("v6")

# Using Aggregate Functions
descriptive_df = spark.sql(
    """
    SELECT POI_ID, round(avg(`Shortest Distance`),3) AS Average_Distance,
    round(stddev(`Shortest Distance`),3) AS Stddev_Distance,
    max(`Shortest Distance`) AS `Max_D`,
    count(Order_ID) AS `Order_Count`
    FROM v6
    Group by POI_ID
    """
)
descriptive_df.show()
```

POI_ID	Average_Distance	Stddev_Distance	Max_D	Order_Count
POI-3	14872.704	1310.943	16235.19	19352

Question 7 For each POI, based on the max distance and orders count from #6, calculate the density using the formula:  $\text{density} = \text{orders\_count} / (\pi * (\text{max\_distance})^2)$

In [109...

```
# create view for SQL analytics
descriptive_df.createOrReplaceTempView("v7")

# Casting Density as Decimal Value
density_df = spark.sql(
    """
    SELECT POI_ID, Order_Count, Max_D, CAST(Order_Count/(3.14*pow(Max_D,2)) AS Decimal(25,15)) as Density
    FROM v7
    """
)
density_df.show()
```

POI_ID	Order_Count	Max_D	Density
POI-3	19352	16235.19	0.000023381989235